

Spam classification with Naive Bayes

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from collections import Counter
from sklearn import feature_extraction, model_selection, naive_bayes, metrics, svm
from IPython.display import Image
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

Loading the Dataset


```
In [2]: data = pd.read_csv('spam.csv', encoding='latin-1')
data.head(n=10)
```

Out[2]:


	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN
5	spam	FreeMsg Hey there darling it's been 3 week's n...	NaN	NaN	NaN
6	ham	Even my brother is not like to speak with me. ...	NaN	NaN	NaN
7	ham	As per your request 'Melle Melle (Oru Minnamin...	NaN	NaN	NaN
8	spam	WINNER!! As a valued network customer you have...	NaN	NaN	NaN
9	spam	Had your mobile 11 months or more? U R entitle...	NaN	NaN	NaN

Distribution spam/non-spam plots

```
In [3]: count_Class=pd.value_counts(data["v1"], sort= True)
count_Class.plot(kind= 'bar', color= ["blue", "orange"])
plt.title('Bar chart')
plt.show()
```



```
In [4]: count_Class.plot(kind = 'pie', autopct='%1.0f%%')
plt.title('Pie chart')
plt.ylabel('')
plt.show()
```




Data Analytics

We want to find the frequencies of words in the spam and non-spam messages. The words of the messages will be model features.


We use the function Counter.

```
In [5]: count1 = Counter(" ".join(data[data['v1']=='ham']["v2"]).split()).most_common(20)
df1 = pd.DataFrame.from_dict(count1)
df1 = df1.rename(columns={0: "words in non-spam", 1 : "count"})
count2 = Counter(" ".join(data[data['v1']=='spam']["v2"]).split()).most_common(20)
df2 = pd.DataFrame.from_dict(count2)
df2 = df2.rename(columns={0: "words in spam", 1 : "count_"})
```

```
In [6]: df1.plot.bar(legend = False)
y_pos = np.arange(len(df1["words in non-spam"]))
plt.xticks(y_pos, df1["words in non-spam"])
plt.title('More frequent words in non-spam messages')
plt.xlabel('words')
plt.ylabel('number')
plt.show()
```



```
In [7]: df2.plot.bar(legend = False, color = 'orange')
y_pos = np.arange(len(df2["words in spam"]))
plt.xticks(y_pos, df2["words in spam"])
plt.title('More frequent words in spam messages')
plt.xlabel('words')
plt.ylabel('number')
plt.show()
```



We can see that the majority of frequent words in both classes are stop words such as 'to', 'a', 'or' and so on.

With stop words we refer to the most common words in a language, there is no single, universal list of stop words.

Feature Engineering

Text preprocessing, tokenizing and filtering of stopwords are included in a high level component that is able to build a dictionary of features and transform documents to feature vectors.

We remove the stop words in order to improve the analytics

```
In [8]: f = feature_extraction.text.CountVectorizer(stop_words = 'english')
        X = f.fit_transform(data["v2"])
        np.shape(X)
```

```
Out[8]: (5572, 8404)
```

We have created more than 8400 new features. The new feature j in the row i is equal to 1 if the word w_j appears in the text example i . It is zero if not.

Predictive Analysis

My goal is to predict if a new sms is spam or non-spam. I assume that is much worse misclassify non-spam than misclassify an spam. (I don't want to have false positives)

The reason is because I normally don't check the spam messages.

The two possible situations are:

1. New spam sms in my inbox. (False negative).
OUTCOME: I delete it.
2. New non-spam sms in my spam folder (False positive).
OUTCOME: I probably don't read it.

I prefer the first option!!!

First we transform the variable spam/non-spam into binary variable, then we split our data set in training set and test set.

```
In [9]: data["v1"] = data["v1"].map({'spam':1, 'ham':0})

        # Splitting Data
        X_train, X_test, y_train, y_test = model_selection.train_test_split(X, data['v1'],
        test_size=0.33, random_state=42)
        print([np.shape(X_train), np.shape(X_test)])

        [(3733, 8404), (1839, 8404)]
```

Multinomial naive bayes classifier

We train different bayes models changing the regularization parameter α .

We evaluate the accuracy, recall and precision of the model with the test set.

```
In [10]: list_alpha = np.arange(1/100000, 20, 0.11)
score_train = np.zeros(len(list_alpha))
score_test = np.zeros(len(list_alpha))
recall_test = np.zeros(len(list_alpha))
precision_test= np.zeros(len(list_alpha))
count = 0
for alpha in list_alpha:
    bayes = naive_bayes.MultinomialNB(alpha=alpha)
    bayes.fit(X_train, y_train)
    score_train[count] = bayes.score(X_train, y_train)
    score_test[count]= bayes.score(X_test, y_test)
    recall_test[count] = metrics.recall_score(y_test, bayes.predict(X_test))
    precision_test[count] = metrics.precision_score(y_test, bayes.predict(X_test))
    count = count + 1
```

Let's see the first 10 learning models and their metrics!

```
In [11]: matrix = np.matrix(np.c_[list_alpha, score_train, score_test, recall_test, precisio
n_test])
models = pd.DataFrame(data = matrix, columns =
                        ['alpha', 'Train Accuracy', 'Test Accuracy', 'Test Recall', 'Test Prec
ision'])
models.head(n=10)
```

Out[11]:

	alpha	Train Accuracy	Test Accuracy	Test Recall	Test Precision
0	0.00001	0.998661	0.974443	0.920635	0.895753
1	0.11001	0.997857	0.976074	0.936508	0.893939
2	0.22001	0.997857	0.977162	0.936508	0.900763
3	0.33001	0.997589	0.977162	0.936508	0.900763
4	0.44001	0.997053	0.977162	0.936508	0.900763
5	0.55001	0.996250	0.976618	0.936508	0.897338
6	0.66001	0.996518	0.976074	0.932540	0.896947
7	0.77001	0.996518	0.976074	0.924603	0.903101
8	0.88001	0.996250	0.976074	0.924603	0.903101
9	0.99001	0.995982	0.976074	0.920635	0.906250

I select the model with the most test precision

```
In [12]: best_index = models['Test Precision'].idxmax()
models.iloc[best_index, :]
```

```
Out[12]: alpha          15.730010
Train Accuracy      0.979641
Test Accuracy       0.969549
Test Recall         0.777778
Test Precision      1.000000
Name: 143, dtype: float64
```

My best model does not produce any false positive, which is our goal.

Let's see if there is more than one model with 100% precision !

```
In [13]: models[models['Test Precision']==1].head(n=5)
```

Out[13]:

	alpha	Train Accuracy	Test Accuracy	Test Recall	Test Precision
143	15.73001	0.979641	0.969549	0.777778	1.0
144	15.84001	0.979641	0.969549	0.777778	1.0
145	15.95001	0.979641	0.969549	0.777778	1.0
146	16.06001	0.979373	0.969549	0.777778	1.0
147	16.17001	0.979373	0.969549	0.777778	1.0

Between these models with the highest possible precision, we are going to select which has more test accuracy.

```
In [14]: best_index = models[models['Test Precision']==1]['Test Accuracy'].idxmax()
         bayes = naive_bayes.MultinomialNB(alpha=list_alpha[best_index])
         bayes.fit(X_train, y_train)
         models.iloc[best_index, :]
```

```
Out[14]: alpha          15.730010
         Train Accuracy  0.979641
         Test Accuracy   0.969549
         Test Recall     0.777778
         Test Precision   1.000000
         Name: 143, dtype: float64
```

Confusion matrix with naive bayes classifier

```
In [15]: m_confusion_test = metrics.confusion_matrix(y_test, bayes.predict(X_test))
         pd.DataFrame(data = m_confusion_test, columns = ['Predicted 0', 'Predicted 1'],
                     index = ['Actual 0', 'Actual 1'])
```

Out[15]:

	Predicted 0	Predicted 1
Actual 0	1587	0
Actual 1	56	196

Conclusion

- We misclassify 56 spam messages as non-spam emails whereas we don't misclassify any non-spam message.

No exercise in this week, but please submit your written file to dropbox via link below:

- <https://www.dropbox.com/request/LnAkVUYXzxAw3dfj1vFK>